



# CS 3110

## Concurrent Programming

Prof. Clarkson

Fall 2015

Today's music:

*Bad Romance* by Lady Gaga;

*Lady Gaga Fugue* by Giovanni Dettori

# Fugue

1. [music] a musical composition in which one or two themes are repeated or imitated by successively entering voices and contrapuntally developed in a continuous interweaving of the voice parts
2. [psychiatry] a disturbed state of consciousness in which the one affected seems to perform acts in full awareness but upon recovery cannot recollect the acts performed

# Review

## Previously in 3110

- Imperative programming: refs, mutable fields, arrays, loops

## Today:

Concurrent programming with two different **libraries**:

- Threads
- Async

# Concurrency

- Networks have multiple computers
- Computers have multiple processors
- Processors have multiple cores

...all working semi-independently

...all sharing resources

**concurrent:** overlapping in duration

**sequential:** non-overlapping in duration

**parallel:** happening at the same time

# Concurrency

At any given time, my laptop is...

- Streaming music
- Running a web server
- Syncing with web services
- Scanning for viruses
- Running OCaml

The OS plays a big role in making it look like those all happen simultaneously

# Concurrency

Applications might also want concurrency:

- Web server that handles many clients at once
- Scientific calculations that exploit parallel architecture to get speedup
- Simulations that model physical processes
- GUIs that want to respond to users while doing computation (e.g., rendering) in the background

# Programming models for concurrency

**Threads:** sequential code for computation

e.g., Pthreads, OpenMP, java.lang.Thread, OCaml **Thread**

**Futures:** values that are maybe not yet computed

e.g., .NET async/await, Clojure, Scala,  
java.util.concurrent.Future, OCaml **Async**

**(and many others)**

**THREADS**



# Threads

- *process*: basic unit of execution; a running program
  - OS manages resources (e.g., memory, privileges) available to process
- *thread*: basic unit of CPU utilization
  - aka *lightweight process* or *task*
  - every process has  $\geq 1$  thread
  - allocation of CPU to threads managed by OS *scheduler*
  - shares most resources with all other threads in the OS process in which it runs
  - but has its own registers and stack (call frames, stack pointer)
- can see these with Mac OS Activity Monitor, Windows Task Manager, Linux **top**

# Threads: concurrent print v1.0

```
let printn n =  
  for i=1 to 1000 do  
    print_int n; print_newline();  
  done
```

```
let t1 = Thread.create printn 1  
let t2 = Thread.create printn 2  
let t3 = Thread.create printn 3
```

To compile:

```
$ cs3110 compile -t file.ml
```

**Problem:** no output...

OCaml starts these threads when it runs, but it doesn't wait for the Threads to finish, so you get 0 output. You need to wait for the Threads to finish to get an output

# Threads: concurrent print v2.0

...

```
let _ = (Thread.join t1,  
         Thread.join t2,  
         Thread.join t3)
```

val join : t -> unit

join th suspends the execution of the calling thread until the thread th has terminated.

**Problem:** output jumbled...

# Threads: concurrent print 3.0

```
let m = Mutex.create ()
```

```
let printn n =  
  for i=1 to 1000 do  
    Mutex.lock m;  
    print_int n; print_newline();  
    Mutex.unlock m  
  done
```

```
let t1 = Thread.create printn 1
```

```
let t2 = Thread.create printn 2
```

```
let t3 = Thread.create printn 3
```

```
let _ = (Thread.join t1, Thread.join t2, Thread.join t3)
```

# Outputs can change

- **Deterministic:** given same input, always produces same output; a mathematical function
- **Nondeterministic:** given same input, sometimes produce different outputs; a mathematical relation

In previous code, nondeterminism from order in which **scheduler** chooses to switch between threads.

# Thread

```
module Thread : sig
  type t
  val create : ('a -> 'b) -> 'a -> t
  val join    : t -> unit
  ...
end
```

- Creating a thread with **create func arg** causes **func arg** to be executed concurrently with other threads
- Scheduler **preempts** threads, executes others, to give **illusion** of parallelism
- Joining a thread with **join t** suspends the calling thread until thread **t** has finished executing

# Mutex

```
module Mutex : sig
  type t
  create  : unit -> t
  lock    : t -> unit
  unlock  : t -> unit
  ...
end
```

- Locking a mutex suspends the calling thread until the lock is available, then acquires the lock and prevents any other thread from locking it
- Unlocking a mutex releases it and allows another (suspended) thread to acquire the lock

**Invariant:** at most one thread can hold the lock at any time

# Other synchronization modules

- **Condition:** enable threads to suspend until some condition becomes true, and to have exactly one thread resume at that time
- **Event:** enable threads to pass messages to one another, and to suspend while waiting to send/receive messages



# Parallelism

- No true parallelism in **Thread** implementation
  - Same as Node.js, ...
  - But not typically true of OS-level threads
- So **speedup** of CPU-intensive computations is not the purpose of this kind of concurrency
- Rather, **latency hiding** of IO-intensive computations is the purpose
  - e.g., while one thread of web server is waiting for file/network read to complete...another thread makes progress on processing requests
  - e.g., while one thread of GUI is calculating new cell in spreadsheet...another thread responds to mouse click

**FUTURES**

# Futures

- *Future*: computation that will produce a value sometime in the future
  - aka *promises* or *delays*
- Completion of computation could be *implicit* (when used, computation forced to occur) or *explicit* (call a function to force computation)
- Computation could be *eager* (starts right away) or *lazy* (starts only when needed)

# Async

- A third-party library for futures in OCaml
- Instead of "futures" calls the abstraction **deferreds**, as in *values whose completed computation has been deferred until the future (and in fact is happening already)*
- Typical use of library is to do asynchronous I/O
  - Launch an I/O operation as a deferred
  - Later on its results will be available
  - Enables *latency hiding*: have multiple I/O operations occurring in parallel

# (A)synchronous I/O

- *Synchronous* aka *blocking* I/O:
  - call I/O function which *blocks*, wait for completion...
  - then continue your computation
  - e.g., `Pervasives.input_line : in_channel -> string`
- *Asynchronous* aka *non-blocking* I/O:
  - call I/O function which is *non-blocking*, function immediately returns, continue your computation, later...
  - I/O completes
  - e.g., `Async.Std.Reader.file_contents : string -> string Deferred.t`
  - how does program make use of completed I/O? ...

# Async: Print file length

```
open Async.Std
```

```
let printlen s =  
  Printf.printf "%i\n" (String.length s)
```

```
let r = Reader.file_contents Sys.argv.(1)
```

```
let _ = upon r (fun s -> printlen s; ignore(exit 0))
```

```
let _ = Scheduler.go()
```

To compile: **cs3110 compile -t -p async filename.ml**

# Scheduler

- Scheduler runs **callbacks** that have been registered to consume the values of deferreds
- Only ever one callback running at a time
  - Async is "single threaded"    **only one callback is running at one time**
  - Like **Thread**, there is no true parallelism, really designed for latency hiding not parallel speedup
- Scheduler:
  - selects a callback whose input has become ready to consume
  - runs the callback with that input
  - never interrupts the callback
    - if callback never returns, scheduler never gets to run again!
    - **cooperative** concurrency
  - repeats

# Deferred so far

```
module Async.Std : sig
  val upon : 'a Deferred.t -> ('a -> unit) -> unit

  module Deferred : sig
    type 'a t
    ...
  end

  module Reader : sig
    val file_contents : string -> string Deferred.t
    ...
  end

  ...
end
```



# Deferred



An `' a Deferred.t` is like a box:

- It starts out empty
- At some point in the future, it could be filled with a value of type `' a`
- Once it's filled, the box's contents can never be changed ("write once")

Terminology:

- "box is filled" = "deferred is determined"
- "box is empty" = "deferred is undetermined"

# Manipulating boxes



- **peek** : 'a Deferred.t -> 'a option
  - use to see whether box has been filled yet
  - returns immediately with **None** if nothing in box
  - returns immediately with **Some a** if **a** is in box
- **upon** : 'a Deferred.t -> ('a -> unit) -> unit
  - use to schedule a callback (the function of type 'a -> unit) to run sometime after box is filled
  - **upon** returns immediately with ( ) no matter what
  - sometime after box is filled (if ever), scheduler runs callback on contents of box
  - callback produces ( ) as return value, but never returned to anywhere

# Creating boxes



- **file\_contents : string -> string Deferred.t**
  - use to read entire contents of file into a string
  - **file\_contents** returns immediately with an empty deferred
  - program can now continue with doing other things (scheduling other I/O, processing completed I/O, etc.)
  - at some point in the future, when file read completes (if ever), that deferred becomes determined
  - any callbacks registered for the deferred will then (eventually) be executed with the deferred
- **return : 'a -> 'a Deferred.t**
  - use to create a deferred that is already determined
- **after : Core.Std.Time.Span.t -> unit Deferred.t**
  - use to create a deferred that becomes determined sometime after a given length of time
  - e.g., **Core.Std.Time.Span.of\_int\_sec 10** represents 10 seconds

# Sequencing boxes



- **bind** :
  - 'a Deferred.t
  - > ('a -> 'b Deferred.t)
  - > 'b Deferred.t
- use to schedule another deferred computation after an existing one
- takes two inputs: a deferred **d**, and callback **c**
- **bind d c** immediately returns with a new deferred **d'**
- sometime after **d** is determined (if ever), scheduler runs **c** on contents of **d**
- **c** produces a new deferred, which if it ever becomes determined, also causes **d'** to be determined with same value

# Sequencing boxes



```
Deferred.bind  
(return 42)  
(fun n -> return (n+1))
```

- first argument is a deferred that is determined with value **42**
- second argument is a callback that takes an integer **n** and returns a deferred that is determined with value **n+1**
- **bind** immediately returns with an undetermined deferred **ud**
- scheduler, when it next gets to run, can notice that first argument is determined, and run callback
- callback gets **42** out of box, **binds** it to **n**, and returns a new deferred that is determined with value **43**
- scheduler can notice that output of callback has become determined, and make **ud** determined with same value

# Sequencing boxes



- (**>>=**)
  - infix operator version of **bind**
  - **bind d c** is the same as **d >>= c**
  - enables a pleasant syntax:
    - **d >>= fun x -> e**
    - should be read as **let x = (contents of) d in e**

# Sequencing boxes

```
Deferred.bind  
(return 42)  
(fun n -> return (n+1))
```

vs.

```
return 42 >>= fun n ->  
return (n+1)
```

# Sequencing boxes

```
open Async.Std
```

```
let sec n = Core.Std.Time.Span.of_int_sec n
```

```
let return_after v delay =  
  after (sec delay)  
  >>= fun () ->  
    return v
```

```
let _ =  
  (return_after "First timer elapsed\n" 5)  
  >>= fun s -> print_string s;  
  (return_after "Second timer elapsed\n" 3)  
  >>= fun s -> print_string s;  
  exit 0
```

```
let _ = print_string "Hello\n"
```

```
let _ = Scheduler.go ()
```



# Recap: Two libraries for concurrency

**Threads:** sequential code for computation

- e.g., Pthreads, OpenMP, java.lang.Thread, OCaml **Thread**

**Futures:** values that are maybe not yet computed

- e.g., .NET async/await, Clojure, Scala, java.util.concurrent.Future, OCaml **Async**

# Upcoming events

- [next Thursday] A4 due

*This is concurrent.*

**THIS IS 3110**